

CS565 - CLOUD COMPUTING

Name : P.V. Srikanth
Roll : 1801CS37

- 1) D - All
- 2) A - Distributed system, Message Passing
- 3) B - False
- 4) A - Utility Computing
- 5) D - Public
- 6) D - All
- 7) D - All
- 8) D - All
- 9) D - All
- 10) D - All of the above

II) Problems with Trap and Emulate process virtualization

(i) Guest OS may realize it is running at lower privilege level.

- * Some registers in X86 reflect CPU privilege level
- * Guest OS can read these values and get offended

(ii) Some X86 instructions which change hardware state run in both privileged and unprivileged modes

- * Will behave differently when guest OS is in ring 0

vs in less privileged ring 1.

- * OS behaves incorrectly in ring 1 will not trap to VMM

- * Instruction set of X86 is not easily visualizable.

(iii) Eflags register is a set of CPU flags. Consider popf

instruction in X86

- * Executed in ring 0, all flags set normally

- * Executed in ring 1, only some flags set

So, popf is a sensitive instruction, not privileged, doesn't trap behaves differently when executed in diff privilege levels, Guest OS is buggy in ring 1.

- * U-rings, no root/non-root nodes yet

- * hypervisor in ring 0, guest OS in ring 1

- * HV doesn't know, so it doesn't try to change setting

- * OS doesn't know, so assumes change was successful.

Binary Translation

In binary translation, we rewrite VM binary to never issue those 17 instructions which causes problems in Trap-and-Emulate.

Goal: Full virtualization

= Guest OS
not modified

Approach: Dynamic binary translation

1. Inspect code blocks to be executed
2. If needed, translate to alternate instruction sequence
Eg: To emulate desired behaviour, possibly even avoiding trap.
3. Otherwise, run at hardware speeds.
cache translated blocks to amortize translation ^{cost}.

12)

There are three key models for device virtualization

- a) Pass-through model: VMM level drivers configures device access permissions
- * VM provided with exclusive access to the device.
 - * VM can directly access the device.
 - * VMM must have exact type of device as what VM expects.

b) Hypervisor-Direct model: VMM intercepts all device accesses

- * Emulate device operation
- * Translate to generic I/O operation.
 - a) Translate to generic I/O operation.
 - b) Transverse VMM-resident I/O state.
 - c) Invoke VMM -resident driver

c) Split-Device Driver model: Device access control split between

- * Front end driver in guest VM
- * Back end driver in Service VM.

13) Linux Containers

A linux container is a set of 1 or more processes that are isolated from the rest of the system. Containers run in host system kernel, separated with policies and can use apps in host system.

- * Linux containers are portable between OS variants supporting linux container and there is no overhead with hypervisor and guest OS kernel.

Advantages of using Containerized Apps over VM-based

- * Containers are independent hosts for applications that use a single, stripped-down version of OS to run whereas VMs use full version of an OS.
- * Containers run a virtualized workload, processed by an application broken up into microservices, making them more lightweight and flexible than a VM.
- * Containers can scale up and down an application quickly and easily.

- * Containers are highly suitable for a microservices architecture in which applications are broken into small, self-sufficient components which can be deployed and scaled individually
- * Containers are an attractive option for deploying and scaling each of these microservices
- * Containers are easily controlled by API and thus are also ideal for automation and CI/CD pipelines

14)

Hotspot detection

Sandpiper implements a hotspot detection algorithm that determines when to migrate VMs. The hotspot detection component employs a monitoring and profiling engine that gathers usage statistics on various virtual and physical services and constructs profiles of resource usage. These profiles are used in conjunction with prediction techniques to detect hotspots in the system. When there is increased workload on a service, it should be simply migrated to a less loaded server. Manually initiated migration lacks the agility to respond to sudden workload changes. So, detecting hotspots is an important step to flexibly remap physical resources to virtual servers.

Hotspot migration:

Upon detecting hotspots, Sandpiper's migration manager is invoked for hotspot mitigation. This hotspot mitigation algorithm resorts to a heuristic to determine which overloaded VMs to migrate and where such that migration overhead is minimized.

Q) Given,

(i) Desired mean response time, $d = 1.5$ s

(ii) The 10 most recent service times (in seconds)

are $1.1, 0.9, 1.2, 0.8, 1.3, 0.7, 1.4, 0.6, 1.2, 0.8$

(iii) Inter-arrival time follows gamma distribution,

$$I(a, b) = I(2, 1)$$

(iv) Peak request arrival time = 1 request per second.
(λ_{peak})

To find the factor by which CPU capacity should be scaled, we need to find the request arrival rate

(λ_{cap}) using the following formula.

$$\lambda_{\text{cap}} \geq \frac{1}{s + \frac{\sigma_a^2 + \sigma_b^2}{2(d-s)}}$$

where S is the mean service time.

σ_a^2 is variance of inter arrival time

σ_b^2 is variance of inter service time

Mean service time (S) = { mean of 10 recent service times}

$$= \frac{1.1 + 0.9 + 1.2 + 0.8 + 1.3 + 0.7 + 1.4 + 0.6 + 1.2 + 0.8}{10}$$

$$= [1 \text{ sec}]$$

Variance of inter arrival time σ_a^2

$$= ab^2 = 2 \times 1^2 = 2$$

(\because Variance of $I(a, b) = a b^2$)

Variance of service time σ_b^2

= Mean of squares } = { square of mean
of service time } of service time

$$= 1.068 - 1.0 = 0.068$$

$$\Rightarrow [\sigma_a^2 = 2] \quad [\sigma_b^2 = 0.068]$$

$$[S = 1] \quad [d = 1.5]$$

$\therefore \lambda_{cap}$ follows the inequality,

$$\lambda_{cap} \geq \frac{1}{1 + \frac{2+0.068}{2(1.5-1)}} = \frac{1}{3.068}$$

$$\boxed{\lambda_{cap} \geq \frac{1}{3.068}} - \textcircled{1}$$

Factor by which CPU capacity should be increased

$$\text{is } \frac{\lambda_{peak}}{\lambda_{cap}} = \frac{1}{\lambda_{cap}}$$

$$\text{As } \lambda_{cap} \geq \frac{1}{3.068}, \quad \underline{\frac{1}{\lambda_{cap}}} \leq 3.068$$

so, the current CPU capacity of the server should be scaled by a max factor of 3.068 so as to service peak request arrival time of 1 request per sec.

16) Microservices architecture allows developers to create separate components of an application through building an application from a combination of small services.

Benefits of microservice based architecture

1) Isolation

Microservices are profitable due to their isolation and resilience. If one of the component fails, developers have the option to use another service and application will run independently.

2) Scalability

It's easier to scale up or down following the requirement of specific application (element).

3) Productivity

Ability to easily understand when compared to monolith. Beneficial to expand your development.

4) Flexibility and Faster project development

Since services work independently, no need to change whole codebase. You can change one component and then deploy. Helps to choose right tool w.r.t task, without affecting communication.

5) Evolutionary

Perfect choice for developers who are unable to predict the kinds of devices the app is going to run on.

17) Factors that should be considered while choosing a right cloud service provider

(i) Regular Compliance → Your business processes running on cloud based system are compliant with the standards which your customers require.

(ii) Adaptable to the Infrastructure requirements → Provider should understand your infrastructure requirements, and offer solutions corresponding to that without compromising on business continuity.

(iii) Favourable service level agreements → Establishes relationship between a user and cloud service provider.

(iv) Different types of cloud services → Selecting the service that is compatible to the business requirement.

(v) Security → Cloud service should ensure that your data is safe and can traverse across different platform resilient with AUTHENTICATION and ease of availability.

- (vi) Won't refuse custom support → Customer service is required
- (vii) Flexibility in pricing plans → Costs should also be considered

18) Black Box monitoring

Black box monitoring refers to the monitoring of services with a focus on areas such as disk space, CPU usage, memory usage.

Grey Box monitoring

Grey box monitoring implies performing analysis of the system at runtime, which brings new limitations to monitorability (The feasibility of solving the monitoring problem).

19) Two approaches to address the problem of networking VMs

- a) One, using specialized Hardware
- b) Software based approach

<u>Hardware</u> <u>based</u>	<u>Software</u> <u>based</u>
(i) SR-IOV, single root I/O virtualization	(i) Open vSwitch
(ii) Sacrifices flexibility and forwarding logic	(ii) Focuses on forwarding flexibility

Hardware based

- (iii) Hits native performance
- (iv) The main idea behind is that CPUs are not designed to forward packets
- (v) SR-IOV, the physical link itself supports virtualization in hardware
- (vi) The NIC provides a physical function (a standard ethernet port). Several virtual functions are also provided
- (vii) Each VM is mapped to one of these functions. So, the VMs themselves get NIC hardware resources

Software based

- (i) The compromise made is to avoid targeting worst case performance.
- (ii) Smart of this approach is the switch routing decisions lie in user space
- (iii) This is where one decides what rules or filters apply to packets of a certain type.
- (iv) Perhaps .based on network updates from other , possibly virtual , such as in network .
- (v) Programmed using open flow .

20) Stages of live VM migration

VM running normally on Host-A

Stage 0 : Pre-migration

Active VM on host A. Alternate physical host may be pre-selected for migration. Block devices mirrored and free resources maintained.

Overhead due to copying

Stage 1 : Pre-reservation

Initialize a container of target host

Stage 2: Iterative pre-copy

Enables shadow paging. Copy dirty pages in successive rounds

Stage 3: Stop and Copy

Suspend VM on host A. Generate ARP to redirect traffic to host B. Synchronize all remaining VM state to host B

Stage 4: Commitment

VM state on Host-A is released

Stage 5: Activation

VM starts on Host-B

Connects to local devices resumes normal operation.

VM running normally on

Host-B

21)

(Given,

Start up that wishes to run a service with -

Cloud service 128 servers (1024 cores)

524 TB storage

Costs are given as follows:

$$\text{cloud service provider cost} = \$0.12 / \text{GB per month}$$

$$= \$0.10 / \text{CPU per hour}$$

$$\begin{aligned}\text{Owning cost} &= \$349 \text{ K for storage} \\ &= \$ \frac{7.5}{\cancel{1024}} \text{ K for CPU} \\ &\quad \text{per month}\end{aligned}$$

Now, we need to compare per month costs of these two options,
 Ownership (Considering operations for M months)

$$\text{Storage} \rightarrow \frac{349}{M}$$

$$\text{Total} \rightarrow \frac{1555 + 7.5}{M}$$

(Considering 0.45 : 0.4 : 0.15 split for hardware : power : network and 3 year lifetime of hardware)

Outsourcing:

$$\text{Storage} = \$0.12 \times 524 \times 1000 \approx \$\underline{\underline{62K}}$$

$$\begin{aligned}\text{Total} &= \text{storage} + \text{CPUs} = \$62K + \$0.10 \times 1024 \times 24 \approx \\ &\approx \$\underline{\underline{136K}}\end{aligned}$$

Breakeven analysis

Now, we shall analyse the required number of months for ownership operation to break even with outsourcing operations.

$$\frac{\$349K}{M} < \$62 \Rightarrow M > 5.55 \text{ months (For storage)}$$

$$\frac{\$1555K}{M} + 7.5 < \$136K \Rightarrow M \geq 12 \text{ months (For total)}$$

- ∴ G would prefer outsourcing since it is taking more than a year for break even point.

22)

Given service period = 30 days

Maximum service hours per day = 15 hours

$$\begin{aligned}\text{Total service uptime} &= 30 \text{ days} \times 15 \text{ hours/day} \\ &= 450 \text{ hours}\end{aligned}$$

Cost per day = ₹ 2500 per day

$$\text{Total cost} = ₹ 2500 \times 30 = ₹ 75000$$

$$\begin{aligned}\text{Total outage time} &= 5 + 0.5 + 1.5 + 5\text{min} + 2.5 \\ &= 9 \text{ hrs } 35 \text{ min}\end{aligned}$$

$$\begin{aligned}\text{Service availability} &= 1 - \frac{\text{outage time}}{\text{uptime(actual)}} = 1 - \frac{9\text{hr } 35\text{min}}{450 \text{ hrs}} \\ &= 97.92\%\end{aligned}$$

So, the monthly uptime (availability) = 97.92% < 99%

$$\begin{aligned}\text{Service credit available} &= 25\% \text{ of total cost} \\ &= \frac{1}{4} \times 75000 = ₹ 18750\end{aligned}$$

$$\therefore \text{Effective cost payable} = 75000 - 18750 \\ = \boxed{₹ 56250}$$

23)

a) Cost / effective hourIn-house server :

Purchase cost = ₹ 70000

$$\text{Purchase cost (per hour)} = \frac{70000}{3 \times 365 \times 24} = \underline{\underline{\text{₹ 2.66}}}$$

Efficiency = 40%

$$\textcircled{a}) \text{ Cost per effective hour} = \frac{2.65}{0.4} = \boxed{6.65}$$

Power &
cooling
↓

$$\therefore \text{Total cost per effective hour} = 6.65 + 20 + 30 \\ P \\ = \boxed{₹ 56.65}$$

Management
cost

Cloud server :

Given cost / hr = ₹ 7

$$\text{Efficiency} = 80\% \quad \textcircled{b}) \text{ Cost per eff hr} = \frac{7}{0.8} = \underline{\underline{\text{₹ 8.75}}}$$

$$\text{The cost per effective hour} = 8.75 + \\ P \\ \text{Management} \\ \text{cost} \\ = \boxed{9.75}$$

$\therefore \text{Total cost / effective hour}$ $\text{cloud} = ₹ 9.75$ $\text{In-house} = ₹ 56.65$

b)

Cost analysis for 3-years:

In-house server:

Purchase cost = ₹ 70000

$$\begin{aligned}\text{Power and cooling cost for 3-years} &= 30 \times 3 \times 365 \times 24 \\ &= ₹ 7,88,400\end{aligned}$$

$$\text{Management cost} = 20 \times 3 \times 365 \times 24 = ₹ 5,25,600$$

$$\text{Total cost} = 1384000$$

Cloud Server:

Cost per hour = ₹ 7

$$\text{Cost for 3-years} = 7 \times 3 \times 365 \times 24 = 183960$$

$$\text{Management cost} = 1 \times 3 \times 365 \times 24 = 26280$$

$$\text{Total cost} = 183960 + 26280 = 210240$$

Given Revenue per day = ₹ 3000

$$\text{For 3 years} = 3000 \times 3 \times 365 = 3285000$$

Expected profits =

$$\text{In-house} = 3285000 - 1384000 = \boxed{19,01,000}$$

$$\text{Cloud} = 3285000 - 210240 = \boxed{30,74,760}$$

c) Let ' γ ' be modified efficiency
Equating cost/effective-hour for both scenarios

Equating cost/effective-hour for cloud = 8.75

From (a), cost/effective-hour for inhouse = 70000 / $3 \times 365 \times 24 \times \gamma$

Cost/effective hour for inhouse = $\frac{70000}{3 \times 365 \times 24 \times \gamma}$

$$\Rightarrow \frac{7}{0.8} = \frac{70000}{3 \times 365 \times 24 \times \gamma} \Rightarrow \gamma = \frac{2.66}{8.75} = 0.304$$

$$\Rightarrow \boxed{\gamma = 30.4\%}$$

\therefore The required efficiency of inhouse = 30.4%

24)

- a) Power consumption in a physical machine is given by,

$$\Delta P_j = \begin{cases} P_{j,\text{idle}} + (P_{j,\text{peak}} - P_{j,\text{idle}}) \times UT_j; & UT_j > 0 \\ 0 & ; \text{otherwise} \end{cases}$$

where,

$P_{j,\text{idle}}$ → Half of peak power consumption $P_{j,\text{peak}}$

UT_j → Percentage of utilization in PM_j

Servers' total energy consumption ΔE_j is determined via the energy model represented as;

$$\Delta E_j = \int \Delta P(UT_j) dt$$

- b) Proposed algorithm first distinguishes the PMs having hotspots, then selects some VMs from each hotspot PM to perform migration and determines destination PMs for the VMs. ~~The detected catastrophes~~

Algorithm to predict hotspots:

Input : PM_{active}

Output : Is PM having hotspot or not

window Length $\leftarrow 20$; $UT_{pm}(t) = (U_{pm}(t)) / (Cap_{pm}(t))$,

Set Th_{hot} applying $MAD = 1 - S^* MAD$

for each $pm \in PM_{active}$ do

begin

if $Length \cdot UTHistory_{pm}(t) < window\ Length$ then

if $(UT_{pm}(t) > Th_{static})$ then

return true

else

return False

else

Predict $UT_{pm}(t+1)$ using CBS Bi-LSTM

if $(UT_{pm}(t) > Th_{hot}) \&\& (UT_{pm}(t+1) > Th_{hot})$ then

return true

else

return False

end

Algorithm to predict cold spots:

We follow same algorithm as hotspot; but we

compare UT_{pm} with Th_{cold}

i.e.) $UT_{pm} > Th_{hot}$ to $UT_{pm} \leq Th_{cold}$.

c) PM's energy consumption relies on the utilization of resources mainly memory, network bandwidth, CPU, etc. However, existing research reveal that the CPU utilizes more power than other resources. Hence, we represent the PM's resource usage by its CPU usage. As compared to the energy consumption caused by different algorithms, the suggested approach uses a lesser amount of energy. The reason being that the PMs having their usage lower than the threshold limit are switched to power saving mode. The energy consumption of PMs is further optimized via packing the VMs into most loaded PMs.